

Does Google de-rank AI-assisted content?

Prepared for BigOrange Marketing following the September 3 pilot review call.

Short answer: No supporting evidence was found in the reviewed sources for a ranking benefit merely because Gemini helped produce content. Current public Google and Bing guidance reviewed here does not document an automatic penalty triggered solely by AI assistance. The defensible risks are low value, duplication, weak oversight, inaccuracy and content produced primarily to manipulate rankings.

What the public guidance supports

Google's published guidance focuses on the purpose and quality of the resulting content, not a preferred production tool. It also warns that generating many pages without adding value can violate its scaled content abuse policy. Bing's current webmaster guidance similarly focuses on large scale publishing without adequate oversight, quality control or editorial review.

That is not a promise that AI-assisted pages will rank or avoid quality systems. It means tool use alone is not a sound release rule.

What the reviewed studies show

Observation	Evidence	Limit
AI detector score had almost no correlation with rank position.	Ahrefs reported $r = 0.011$ across 600,000 pages in July 2025.	Observational and detector dependent; it does not establish causation.

Observation	Evidence	Limit
High AI share pages received fewer impressions in a separate sample.	Ahrefs reported roughly two to three times fewer impressions in a Search Console sample of 80,861 pages in July 2026.	The authors identify authority, content type and recency as confounds.
Human classified pages held the number one result more often in one study.	Semrush reported 80 percent versus 9 percent across 42,000 pages and 20,000 keywords in April 2026.	AI detectors are imperfect and the comparison is not causal.
One new domain cohort lost visibility after early gains.	SE Ranking followed 2,000 AI generated articles across 20 new domains. Around the three month mark, the share appearing in the top 100 fell from 28 percent to 3 percent and remained low during the 16 month observation window.	Publisher run longitudinal cohort without a matched control group. It does not establish that AI use caused the decline.

What this means for the builder authority system

BigOrange is not directly comparable to a set of brand new domains. The September 1, 2026 public crawl recorded 299 published posts. Emelia's Moz export dated August 14, 2026 showed 17 tracked builder terms on the existing builder pillar, including 13 in positions 1 through 3. That is a dated snapshot, not a claim about current rankings.

The practical release risk is a cluster of closely related pages without distinct reader jobs, approved expertise, editorial review or staged quality assurance. Local schema candidates and answer blocks do not prove usefulness or publication readiness.

Recommended release controls

1. Give every approved page one distinct reader job and accurate visible content.
2. Add Janice's expertise only where factual and public use permission is recorded.
3. Keep a named human editor accountable for each page.
4. Stage releases and evaluate each page before expanding the set.
5. Treat Article JSON-LD as partial until author, dates, image, final URL, visible content, plugin ownership and rendered validation are approved.

Correction on the record

The claim that adding Gemini output helps rankings is withdrawn. No supporting evidence was found in the reviewed sources. The safe decision is based on content quality, distinct purpose, approved expertise and human review, not the name of the drafting tool.

Sources reviewed

- [Google guidance about AI generated content](#)
- [Google spam policies](#)
- [Google generative AI content guidance](#)
- [Bing Webmaster Guidelines](#)
- [Ahrefs correlation study, July 7, 2025](#)
- [Ahrefs search performance study, July 27, 2026](#)
- [Semrush ranking study, April 1, 2026](#)
- [SE Ranking longitudinal cohort report, March 19, 2026](#)

Evidence boundary: the large studies are not peer reviewed causal evidence. They use imperfect detectors, observational comparisons, or publisher run cohorts without matched controls.