

Deborah Mara

ANSWER ENGINE OPTIMIZATION

Getting found when the search box is a *conversation*.

A configuration-first plan to make Deborah the agent that ChatGPT, Perplexity, Google AI Overviews and Claude name when someone asks who to trust in Monmouth, Ocean and Mercer County — without writing a single blog post.

THE SHIFT

Buyers stopped clicking. They started asking.

"Who's a good realtor in Monmouth County?" used to return ten blue links. Increasingly it returns one paragraph and three names. If Deborah is not one of the three, the other seven links no longer matter.

Answer engines do not rank pages the way search engines do. They assemble an answer from sources they can **read, verify and corroborate**. That changes what wins. Volume of content matters less. Being unambiguously machine-readable matters enormously.

WHY THIS SUITS DEBORAH SPECIFICALLY

Deborah is not going to write blogs, and this plan does not ask her to. Almost everything that drives AI visibility is **configuration, not composition** — structured data, crawler permissions, entity consistency and a handful of evergreen pages built once. It is a fixed amount of work, not a treadmill.

What the engines already have to work with

Her credibility is unusually strong for a single agent. These are the facts an answer engine would quote — if it could read them.

\$145M+

CLOSED SINCE 2013

13

YEARS WITH RE/MAX

32

TRANSACTIONS IN 2025

100+

FIVE-STAR REVIEWS

THE FINDING THAT STARTS THIS WHOLE PLAN

Her RE/MAX and RealTrends recognition are published on the site as **image files** — `remax-819x1024.jpg` and `realtrends-819x1024.jpg` — with **empty alt attributes**. Every proof point she has earned is currently locked inside a JPEG. To an AI crawler, that recognition does not exist.

Nothing else in this document delivers more for less effort than moving those facts into text and structured data.

PILLAR ONE

Let the AI crawlers in — explicitly.

The site's `robots.txt` today contains a bare `Disallow:` under a single wildcard rule. That is permissive by accident rather than by decision, and it names none of the agents that actually matter now. Make the invitation explicit, and audit it later.

The three OpenAI agents do different jobs and all three are worth allowing: `GPTBot` gathers training data, `OAI-SearchBot` builds the search index that ChatGPT cites, and `ChatGPT-User` fetches a page live when a user's question needs it. Blocking the second two removes her from citations.

```
# Answer engines — allow
User-agent: GPTBot
User-agent: OAI-SearchBot
User-agent: ChatGPT-User
User-agent: PerplexityBot
User-agent: Perplexity-User
User-agent: ClaudeBot
User-agent: Claude-SearchBot
User-agent: Claude-User
User-agent: Google-Extended
User-agent: Applebot-Extended
User-agent: meta-externalagent
Allow: /

# Everything else
User-agent: *
Allow: /

Sitemap: https://[her-domain]/sitemap_index.xml
```

DO THIS BEFORE LAUNCH, NOT AFTER

The staging site is currently `index, follow` with a live sitemap and a page title of **"Home - WordPress Blog"**. Crawlers may already be indexing an unfinished site under a default name. Either `noindex` it until launch, or finish the basics now. Both are settings changes, not development.

PILLAR TWO

Publish an `llms.txt`.

An emerging convention: a single plain-language file at `/llms.txt` that tells a language model what this site is and what matters on it, without making it infer from navigation. It costs one file. Adoption is still uneven, so treat it as cheap upside rather than a guarantee.

```
# Deborah Mara – RE/MAX Homeland REALTORS
```

```
> Real estate agent serving Monmouth, Ocean and Mercer  
> counties, New Jersey. 13 years with RE/MAX, $145M+ in  
> closed transactions since 2013, 32 transactions in 2025,  
> 100+ five-star client reviews. Specialises in active adult  
> and 55+ communities in central New Jersey.
```

```
## Key pages
```

- [Selling in central NJ](/selling): timing, pricing, what to expect
- [What your home is worth](/home-value): how NJ valuation works
- [55+ communities](/active-adult): comparison of age-restricted communities
- [Choosing an agent](/choosing-an-agent): questions to ask before signing

```
## Contact
```

- Phone: 917.747.5055
- Email: marasurrealestate@gmail.com

Structured data — say it in a language machines cannot misread.

Schema.org JSON-LD is how a machine learns, without guessing, that Deborah is a real estate agent, at a named brokerage, serving three specific counties, with a verifiable track record. This is the single highest-leverage item in the plan and it is pure configuration.

```
{
  "@context": "https://schema.org",
  "@type": "RealEstateAgent",
  "name": "Deborah Mara",
  "telephone": "+1-917-747-5055",
  "email": "marasurrealestate@gmail.com",
  "url": "https://[her-domain]/",
  "image": "https://[her-domain]/deborah-mara.jpg",
  "areaServed": [
    {"@type": "AdministrativeArea", "name": "Monmouth County, NJ"},
    {"@type": "AdministrativeArea", "name": "Ocean County, NJ"},
    {"@type": "AdministrativeArea", "name": "Mercer County, NJ"}
  ],
  "worksFor": {
    "@type": "Organization",
    "name": "RE/MAX Homeland REALTORS"
  },
  "knowsAbout": [
    "Active adult communities", "55+ housing",
    "Residential resale", "Downsizing", "Central New Jersey"
  ],
  "sameAs": [
    "https://www.homes.com/real-estate-agents/deborah-mara/1g6qce5/",
    "https://www.ratemyagent.com/real-estate-agent/deborah-mara-b073h2/",
    "[Facebook]", "[Instagram]", "[LinkedIn]", "[Google Business Profile]"
  ]
}
```

THE SAMEAS ARRAY IS DOING THE HEAVY LIFTING

Answer engines resolve identity by **corroboration**. When the same name, phone and brokerage appear consistently across her site, Google Business Profile, homes.com, RateMyAgent and her social profiles, she becomes a confidently-known entity rather than a string of text. That confidence is what earns a citation.

Also add

Schema type	Where and why
FAQPage	On each answer page. The format answer engines most readily lift from — a question, then a direct answer.
Person	Deborah herself, linked to the RealEstateAgent entity via worksFor and jobTitle .
BreadcrumbList	Once real URLs exist. Gives crawlers site structure.
AggregateRating	Only if the 100+ five-star reviews can be tied to a real, verifiable review source. Never fabricate or self-assert a rating — it violates platform policy and is the fastest way to lose trust signals entirely.

PILLAR FOUR

One identity, stated identically, everywhere.

The cheapest win available, and it requires no writing at all. Name, phone, brokerage and service area must match *character for character* across every profile. A mismatch — "Deb" in one place, a different phone in another — splits her into two weaker entities that neither engine fully trusts.

Property	Action
Google Business Profile	Verify categories, service area, hours, phone. This is the strongest local signal there is.
homes.com profile	Already the chosen listings source. Confirm bio, contact and brokerage match.
RE/MAX Homeland roster	Confirm the brokerage page lists her with the same details.
RateMyAgent	Live and already carrying reviews. Link it in <code>sameAs</code> .
Facebook / Instagram / LinkedIn	Same name format, same phone, link back to the owned site.
soldbymara.com	RE/MAX-owned, no publish access. Do not fight it — point the owned site at the same entity so both corroborate.

PILLAR FIVE

Four pages instead of a blog.

Answer engines cite pages that **answer a question completely**. They rarely cite a homepage. Deborah needs four evergreen pages, built once by us, not a posting schedule she will not keep.

Page	The question it owns
Central NJ market reality	"Is now a good time to sell in Monmouth / Ocean / Mercer?"
What your home is worth	"How do I find out what my house is worth without a Zillow guess?"
55+ communities in central NJ	"What are the best active adult communities in central New Jersey?"
Choosing an agent in NJ	"How do I pick an agent, and what should I ask?"

THESE PAGES ARE LOAD-BEARING FOR THE WHOLE PLAN

They are also exactly what ChatGPT Ads requires — its documented best practice is an answer-first, research-stage landing page — and what Google Search needs for quality score, and what Meta needs as a credible destination. **Build once, point four channels at it.**

A structural blocker to fix first

The site is currently one page with anchor links — [#about](#) , [#buysell](#) , [#listings](#) . An answer engine cannot cite a fragment. Without distinct URLs there is nothing for it to point at, and the four pages above cannot exist.

PILLAR SIX

Technical hygiene the crawlers actually check.

Item	Status today
Server-rendered HTML	Good. WordPress serves real HTML. Content is not hidden behind JavaScript, which is where most modern sites lose crawlers.
Site title	Broken. "WordPress Blog" — one untouched field in Settings → General, which also poisons every social share preview.
Heading hierarchy	Broken. No <code><h1></code> anywhere; headings start at <code>h2</code> .
Image alt text	Broken. Empty on every image, including both award badges.
Meta description	Missing. Yoast is installed and can set it.
Internal links	Broken. 19 of 32 links resolve to <code>#</code> or nothing.
Sitemap	Live, but publishing author and category sitemaps for content that does not exist.

MEASUREMENT

How we will know it is working.

AI visibility is harder to measure than rankings, so measure it three ways and do not pretend any one of them is complete.

- Referral traffic in GA4.** Segment referrers for `chatgpt.com`, `perplexity.ai`, `claude.ai` and Gemini. Small numbers, high intent.
- Crawler hits in server logs.** Confirm GPTBot, OAI-SearchBot, PerplexityBot and ClaudeBot are actually fetching pages. This is the earliest signal that the configuration took.
- Prompt testing, monthly.** Ask the engines directly — "best realtor in Monmouth County NJ", "who should I talk to about selling a 55+ home in central New Jersey" — and record whether she is named. Manual, imperfect, and the only true measure of the outcome we are buying.

SET EXPECTATIONS HONESTLY

Answer engines re-crawl and re-index on their own schedules. Structured data and crawler permissions can take **weeks** to surface, and no agency can guarantee a citation. What we control is making her the most machine-legible, best-corroborated agent in her three counties — which is what moves the odds.

SEQUENCE

What to do, in what order.

Ordered by return on effort. The first four items are settings changes and can be done the day we have site access.

#	Action	Effort	Priority
1	Fix the site title and meta description in Settings → General and Yoast	5 minutes	CRITICAL
2	Move award and credibility facts out of JPEGs into real text, and add alt text	1 hour	CRITICAL
3	Add the <code>RealEstateAgent</code> + <code>Person</code> JSON-LD with full <code>sameAs</code>	2 hours	CRITICAL
4	Publish the explicit <code>robots.txt</code> and <code>llms.txt</code>	30 minutes	CRITICAL
5	Add a lead form and tappable phone — nothing converts without it	1 hour	CRITICAL
6	Audit NAP consistency across GBP, homes.com, RE/MAX, RateMyAgent, socials	2 hours	HIGH
7	Build the four answer pages on real URLs with <code>FAQPage</code> schema	1 week	HIGH
8	Replace the Paris stock photo and the placeholder listings	2 hours	HIGH
9	Fix the 19 dead links and promote the headline to <code>h1</code>	1 hour	MEDIUM
10	Stand up AI referral tracking and the monthly prompt test	2 hours	MEDIUM

What we need from Deborah's side

Only one thing blocks all of it: **administrator access to the WordPress site**. Application passwords are already enabled, so adding `dillonmohr8777@gmail.com` as an administrator takes one click from whoever holds access today, and can be revoked just as easily.

ONE CONSTRAINT THAT GOVERNS EVERY PAGE AND AD

US real estate advertising falls under the **Fair Housing Act**. Copy must never suggest a preference or steer by race, colour, religion, sex, familial status, national origin or disability. The 55+ material is lawful under the Housing for Older Persons exception, but only when it refers to genuinely age-restricted communities. Her broker at RE/MAX Homeland REALTORS should review the pages before they go live.

Findings in this document were read from the live staging site on 9 September 2026. Crawler names and the `llms.txt` convention reflect current published practice and should be re-verified at implementation, since the agents and conventions are still changing.
