



The system worked. The model did not earn promotion.

Chronos-2 completed the governed forecast path and beat the trailing-four-week mean. It finished slightly behind persistence, while its uncertainty band covered only 7 of 12 held-out weeks.

RETAIN / EVIDENCE ONLY

Improve the target data and run rolling-origin tests before any planning or automation proposal.

Chronos WAPE

25.75%

held-out evaluation

Persistence WAPE

25.11%

simple baseline won

Band coverage

58.33%

7 of 12 vs intended 80%

History warning

73.39%

in two flagged weeks

What we integrated

VERIFIED PATH · READ ONLY

01

Account verified

The intended Momentum 360 account was confirmed before extraction.

02

Read-only aggregate

Only weekly aggregate CRM evidence entered the pilot.

03

Regular time series

A complete 44-week contacts-created series was formed.

04

Governed gate

Client, model, values, horizon, and request must match.

05

Local model run

Chronos-2 produced the point and quantile outputs.

06

Baseline score

The held-out result was compared with simple baselines.

07

Safe abstention

Persistence won, so the system retained evidence only.

Target boundary


Weekly HubSpot contacts created means CRM creation activity. It does not mean verified leads, qualified leads, conversions, deals, demand, or revenue.



The simplest baseline won.

The model result is useful because the governed path abstained.

— Actual — Chronos p50 - - Persistence □ p10–p90



Twelve-week held-out evaluation of contacts created. The wide p10–p90 band is shown in full; values are displayed as whole contacts.

Held-out error comparison

METHOD	MAE	RMSE	WAPE	SMAPE
Chronos-2	29.91	35.66	25.75%	24.94%
Persistence	29.17	35.11	25.11%	24.40%
Trailing four-week mean	37.58	44.55	32.35%	29.91%

0.64 points

Persistence advantage on WAPE

9 of 12

Weeks where p50 exceeded actual

311 contacts

Mean p10–p90 interval width

Executive read

Chronos beat the trailing mean, but one held-out window does not establish production value. Promotion requires repeatable baseline wins and better-calibrated uncertainty.



Two weeks dominate the history.

Likely import, migration, or backfill activity must be explained before the target can represent organic demand.


73.39%

of complete-window records sit in two weeks

10,895 of 14,846 contacts created. The source evidence remains unchanged; the second chart only removes those two values from its line for legibility.


Full history

Nothing removed



Normal-range zoom

Flagged weeks marked



What is defensible

- 44 complete weekly steps
- No missing weeks on the regular grid
- No invalid or missing created dates
- Read-only aggregate extraction

What remains blocked

- Qualified-lead and conversion forecasting
- Revenue, spend, and channel forecasting
- Conversation forecasting
- Production promotion

Conversion reporting is pending validation.



Improve the target before improving the model.

Build a cleaner signal, then require the model to win repeatedly.

01 Explain the two spikes

Use HubSpot import and audit history to classify the 2025-11-17 and 2026-02-02 concentrations.

02 Define an organic target

Exclude test, spam, duplicate, and imported records without rewriting raw evidence.

03 Capture historical transitions

Preserve immutable source and lifecycle events before channel or qualified-lead forecasting.

04 Add approved future facts

Use scheduled webinars, promotions, holidays, and fixed budget calendars only when documented.

05 Run rolling-origin tests

Require Chronos to beat persistence and a suitable seasonal baseline across several windows.

06 Recalibrate uncertainty

Require observed p10-p90 coverage near the intended 80% before promotion.

PROMOTION GATE

Baseline win across rolling windows

CALIBRATION GATE

Observed p10-p90 coverage near 80%

Verified boundaries

WHAT THE PILOT DID NOT DO

No CRM writes

No spend authority

No sending authority

No conversion claim

No scheduled automation

No production promotion

MODEL

Amazon Chronos-2 · local CPU lane

DATA CLASS

Sanitized aggregate evidence

EVALUATION

32 context weeks · 12 held-out weeks

Why this matters

The architecture prevented a confident-looking forecast from becoming a business instruction. That abstention is the successful result of this pilot.