

AI lead follow-up: the sequence that does not sound like a robot

Dillon Mohr · Momentum Digital

PRIVATE REVIEW EDITION · SEPTEMBER 2026

Complete source draft, formatted for your review. Publication status and open decisions are preserved inside.

Before this goes public

This is an internal review copy. Offers, pricing, first-person claims, client examples and delivery promises require their existing approvals before public use. This PDF export does not change the source approval status.

Source: 07-ai-lead-follow-up-sequence-not-a-robot.md

Typography: Archivo Black display, Georgia reading text and Calibri utility text. Georgia is a print-specific reading choice; the existing Momentum web design system is unchanged.

Inside

01 / Why do most automated sequences sound like a robot?

02 / What does draft-then-human-sends look like?

03 / What is the sequence?

04 / How fast should I follow up?

05 / What should I measure?

06 / What should stay human?

07 / FAQ

08 / The next step

title: "AI lead follow-up: the sequence that does not sound like a robot" slug: ai-lead-follow-up-sequence-not-a-robot target_query: "lead follow up", "ai lead follow up" and "automated lead follow up" intent: informational (commercial CPC) feeds_ebook: "B - From Missed Call to Booked Job" meta_description: "The machine drafts, a person sends. The four-touch sequence, what stays human, and the five columns to measure, from how we run our own inbox." publish_week: 4 sources: - C:\Users\dillo\Documents\Codex\projects\client-operations\clients\momentum-360\deliverables\2026-09-07-google-aistudio-batch\research\aeo-geo-topic-research.md - C:\Users\dillo\Documents\Codex\projects\client-operations\clients\momentum-360\deliverables\2026-09-07-google-aistudio-batch\research\evidence-bank.md - C:\Users\dillo\Documents\Codex\projects\client-operations\clients\momentum-360\deliverables\2026-09-04-ai-division-plan\PLAN.md

AI lead follow-up: the sequence that does not sound like a robot

By Dillon Mohr, Momentum Digital, Philadelphia

AI lead follow-up works when the machine drafts and a person sends. The draft lands in your inbox with the lead's name, what they asked, and a reply written in your words. You read it, fix one line, hit send. Speed comes from the draft. Trust comes from the human. Do that four times, then stop.

That's the system. Here's why it's built that way and how to run it.

Why do most automated sequences sound like a robot?

Three reasons, and none of them is the AI.

They're written once, for everyone. They never mention what the lead actually said. And they're sent by nobody, so nobody's name is on the hook when they land wrong.

The search numbers for this are small and expensive. "Lead follow up" gets about 210 searches per month in the US. "AI lead follow up" is about 40 a month, but advertisers pay close to \$39 a click for it. "Automated lead follow up" is about 20. Low volume, high price means buyers, not browsers. People with a real inbox problem.

I'll give you our own failure so you know I'm not lecturing. In two days in early September 2026 we sent 241 cold emails to businesses in two cities. Two bounced. Two auto-replied. Zero humans wrote back. That was cold outreach, not lead follow-up, but the lesson transfers: nothing in those emails was about the person receiving them, and people can tell.

What does draft-then-human-sends look like?

It's how our own inbox runs today.

Every six hours a drafter reads the new threads in Gmail and Slack, sorts the ones that need a human from the bounces and notifications, and writes one unsent draft per thread in my voice. It never sends. It can't. I open the draft, fix a line, and send it myself, or delete it.

Lead follow-up is the same shape:

1. A lead arrives. Form, call, text, whatever.
2. The machine checks whether it's a duplicate of someone already in the system.
3. It pulls what the lead actually asked for.
4. It drafts a reply in your voice that refers to that request.
5. It assigns an owner by name.
6. The owner sends it.

Step 6 is the whole point. The reply comes from a person because it was sent by a person.

What is the sequence?

Four touches. Then stop.

To, within minutes, automatic. An acknowledgment. "Got your request about [service]. A person will reply during business hours today." This one can be fully automated because it promises nothing a person hasn't agreed to keep.

T1, same business day, human-sent. One specific line about their request, and one question or one time slot. "You mentioned the back deck is about 300 square feet. Is it ground level or raised? I can look Thursday afternoon."

T2, two days later, human-sent. The specific thing. Something you noticed about their job, their address, their timing. Not "just checking in". Never "just checking in".

T3, five days later, human-sent. The close with an out. "Last one from me. If you still want that estimate, reply and I'll set it up. If not, no hard feelings."

Then log the disposition and stop. Four is enough to be helpful. Five is enough to be ignored.

How fast should I follow up?

Here's the honest answer: I don't have a measured baseline I'd publish for our own accounts yet. Everybody in this space quotes somebody's study. I'd rather you measure your own.

Add one column to your lead sheet: time from lead received to first human reply. Fill it in for a month. Then cut it in half. The acknowledgment can be seconds. The human reply should be the same business day. Past that, you're one of four quotes instead of the first call.

What should I measure?

Five columns, and the first two matter most.

1. **Time to first response.** Any response, including the automatic one.
2. **Time to first human response.** The one the lead remembers.
3. **Exceptions per week.** Leads the machine couldn't sort or route. If this number climbs, your intake changed and nobody told the automation.
4. **Duplicates caught.** Every duplicate you don't catch becomes two follow-up sequences to the same person.
5. **Disposition.** Booked, quoted, not a fit, spam, no reply. Whole numbers. No estimates.

Dedup before you draft. For a Pittsburgh painting company we reviewed every automation run over four weeks in summer 2026: 22 runs, 4 real inquiries, 18 solicitations, tests and duplicate replies. The old parser had also misread several names and phone numbers, so a few of those "leads" would have been followed up under the wrong name. We corrected them by hand. If you skip this step, your AI will write a warm, personal follow-up to a spam bot in your own voice.

What should stay human?

Sends. Every one. Prices. Complaints. Anything medical or legal. Anything with a dollar sign in it.

The rule we run our own operations on is simple. Reading and drafting can happen unattended. Sending, posting, spending and deleting need a person's hand. An AI that follows up on leads is a drafter with a good memory. The moment it starts sending on its own, you've handed your reputation to something that can't be embarrassed.

FAQ

Can AI write the follow-up in my voice? Yes, if you write the voice down first: the words you use, the words you never use, how you sign off. Without that file it writes in everybody's voice, which is nobody's.

How many follow-ups is too many? Four, then stop. The fourth gives them an out. Past that you're training people to ignore you.

What's the difference between this and a chatbot? A chatbot talks to the lead. This drafts for you, and you talk to the lead. For a service business the second one builds the relationship you actually want.

The next step

If you'd like a look at how leads reach you and how long they sit, I'll do a 15-minute snapshot: three observations with the evidence for each, and one suggested next step. Free. No deck, no pitch. Reach me through needmomentum.com.

Momentum built the machine. Now we build yours.