

Momentum AI division: launch and scale proposal

Dillon Mohr · Momentum Digital

PRIVATE REVIEW EDITION · SEPTEMBER 2026

Complete source draft, formatted for your review. Publication status and open decisions are preserved inside.

Before this goes public

This is an internal review copy. Offers, pricing, first-person claims, client examples and delivery promises require their existing approvals before public use. This PDF export does not change the source approval status.

Source: PLAN.md

Typography: Archivo Black display, Georgia reading text and Calibri utility text. Georgia is a print-specific reading choice; the existing Momentum web design system is unchanged.

Inside

- 01 / The decision
- 02 / What already exists
- 03 / Initial customer and qualification
- 04 / Three offers with finite delivery contracts
- 05 / Economics and capacity
- 06 / Sales motion and the next seven working days
- 07 / Thirty, sixty and ninety days
- 08 / Ownership and compensation discussion
- 09 / The operating system and model roles
- 10 / Risks and concrete responses
- 11 / Meeting decision sheet
- 12 / The thesis, and whether it survives contact with the evidence
- 13 / Why this changes the plan's centre of gravity
- 14 / Market position — anchors for the pricing decision
- 15 / Mac's open asks for the sales deck
- 16 / Match-back: the differentiator, and the defect underneath it
- 17 / The design proof, with the numbers corrected
- 18 / Control-system gaps that gate the launch
- 19 / Two client threads that bear on the launch
- 20 / Security — exposure state, locators only
- 21 / Open items — none of these are blockers, all of them are decisions
- 22 / What this extension does not change
- 23 / Corrections, 2026-09-07 evening

Momentum AI division: launch and scale proposal

Prepared for Dillon Mohr and Mac Frederick, September 4, 2026. State: reviewed proposal for the planning meeting. Prices, staffing, compensation and sales targets below are proposed assumptions. No offer has been approved or sent through this work.

The decision

Launch a three-client founding cohort around a single promise: **make a service business easier to discover, faster to respond, and simpler to operate, with evidence of what changed.** Start with one flagship lead-operations workflow. Offer AI visibility and campaign production as tightly scoped alternatives for accounts whose needs fit them. Keep Mac's four service categories for navigation; do not launch four separate delivery organizations.

The competitive advantage is Momentum's customer relationships, understanding of marketing, implementation, and accountable support. Model access is an input. A customer can buy a chatbot directly; the division must earn its fee by selecting a valuable process, connecting the right systems, verifying the result, and maintaining it.

The first-month target is three signed paid pilots, with at most one completed and referenceable implementation by day 30, and measured delivery economics by day 30. This is a target, not a sales forecast. Referenceability requires client permission, an accepted pilot, and measured outcome evidence. Do not promise three completed or referenceable implementations in week one. Schedule the remaining implementations against actual capacity. The plan does not depend on building a proprietary foundation model or purchasing hardware. [Correction applied 2026-09-07 per QWEN27-REVIEW.md defect 1.]

What already exists

Mac's September 4 messages explicitly ask Dillon to focus on building the AI division and discuss agency growth. He also mentions Jesse pitching selected packages and possible split commissions. The split, budget and decision rights are unresolved. Sources: [focus/reassignment](#), [division discussion](#).

Mac has named AEO/GEO, AI Design, AI Marketing and AI Automation as service lanes and proposed a free branded audit funnel. The updated agency deck still needs decisions on AEO scope and price, current price anchors, guarantee language and named case-study use. Sources: [service plan](#), [audit thread](#), [deck thread](#).

The evidence packet identifies one signed \$500/month ChatGPT Ads management lane inside a broader \$1,500/month package. That is evidence of a sale, not collected AI-division revenue or proven fulfillment margin. The \$7,400 expected September commission sheet is mixed portfolio income and must not become the division's starting MRR. Existing deck, logo, reporting and routing systems are reusable; the observed 26-slide deck turnaround was about 1h58, not an established under-one-hour SLA.

Full source and limitations: BUSINESS-EVIDENCE.md. That packet uses targeted recent Slack searches, two full threads and current local artifacts; it is not a complete historical audit.

Initial customer and qualification

Recommended first segment: established local service businesses already connected to Momentum, with a named decision-maker, a working website, an existing lead source and CRM or booking process, and a repeatable operational problem. Prefer a low-sensitivity workflow such as form triage, internal response drafting or source reconciliation. Do not use healthcare or legal case records in the founding experiment. The initial segment is a hypothesis grounded in Momentum's delivery base, not a measured market-size claim.

Select accounts by five observable conditions: relationship owner approves the discussion; the buyer describes a repeated costly problem; baseline records exist; necessary system access is obtainable; an employee owns adoption and exceptions. Any missing condition becomes discovery work. An annual-revenue range does not establish willingness to pay.

The exact registry explicitly affiliates Bridge Software, Pritzker Law Group, Nexla and Puttery NYC with Momentum, but their fit differs and affiliation alone does not authorize outreach. Puttery's access dependencies make it an engineering proof opportunity, not the simplest first rollout. Have Mac and the account owner choose three appropriate buyers from the actual sales pipeline. Do not use Dillon's entire 26-client portfolio as a Momentum prospect list.

Three offers with finite delivery contracts

All figures are proposed USD prices, excluding media spend, taxes and third-party pass-through fees. Present a written 30-day pilot scope and renewal terms; do not silently amend existing client contracts. A retainer buys maintenance and specified improvements, not unlimited new builds.

Offer	Proposed setup / monthly	Included scope	Acceptance and support
Lead Operations Pilot — flagship	\$3,500 / \$1,500	One intake source, one destination CRM or sheet, one routing workflow, duplicate protection, internal response draft, exception queue, one dashboard and operator handoff. One approved improvement monthly.	Acceptance is met only when 30 named test cases are executed, with at least 95% passing, all failures documented with owner remediation plan, destination readback succeeds for 100% of accepted test records, no unresolved cross-account write failures remain, and the named operator successfully completes 5 supervised cases. Mac or the named account owner signs the acceptance scorecard before recurring support begins. Up to 2 support hours and one 30-minute review monthly, included within the 6-hour total budget. [Correction applied 2026-09-07 per QWEN27-REVIEW.md defect 4.]

Offer	Proposed setup / monthly	Included scope	Acceptance and support
AI Visibility Program	\$1,500 / \$1,500	One domain/location, baseline of 20 buyer questions across two selected engines, technical/content priorities, two useful page improvements monthly, repeated measurement and source-linked report.	Each reported observation has engine, prompt, date and source. Page changes pass technical and editorial QA. One consolidated revision round, one 30-minute review; total delivery budget 6h/month.
Campaign Production System	\$750 / \$1,250	One brand, one campaign brief, four original copy/design concepts and up to eight straightforward size variants, plus one results/readiness summary. Reuse approved brand assets.	Exact brief and dimensions, brand and claims review, editable agreed files, one consolidated revision round. No shoot, custom animation, paid placement, unlimited video or new website. Total budget 6h/month.

Lead Operations setup assumes 24 total delivery hours; Visibility 10; Campaign 5. These include review and handoff. Measure them. Lead Operations acceptance requires all 30 agreed tests passing, complete readback for accepted records, zero unresolved critical defects, five successful operator cases, and written acceptance by the client decision-maker through the account owner. Recurring support begins at that accepted handoff. Client delays pause the delivery clock; support caps do not excuse defects in the agreed acceptance criteria. Define business-hours acknowledgment before sale; do not promise 24/7 uptime or emergency response without paid coverage.

The visibility program measures observed presence and improves source quality. It cannot guarantee rankings, citations, qualified leads or revenue. Report engine variation and referral evidence separately. Campaign production is a standalone deliverable; creative quantity is not proof of business lift. ChatGPT Ads remains a separate existing service decision, not an assumed feature of these three offers.

Keep the free audit Mac requested: a **15-minute qualification snapshot** with three evidenced observations and a suggested next step. Cap production at 30 minutes after a reusable template exists. Deeper diagnosis is a paid scoped engagement or part of setup. Do not replace the existing free funnel with an arbitrary \$497 audit merely because a model suggested it.

Economics and capacity

economics.py generates the CSV and JSON. Assumptions: \$75/hour loaded delivery labor, 10% sales allowance, \$500/month fixed operating allowance, per-offer tool allowances, full collections. The labor rate and sales allowance are placeholders, not Dillon's compensation. Do not double-count a later agreed salary or delivery commission. Contribution is before taxes, financing, additional management time and any unmodeled costs; it is not net profit.

Tool allowances: \$325/month for the three-client cohort, \$650/month for six clients, and \$1,300/month for twelve clients. Contribution = monthly revenue minus delivery labor at \$75/hour minus 10% sales allowance minus \$500 fixed operating allowance minus tool allowances. [Correction applied 2026-09-07 per QWEN27-REVIEW.md defect 3 — these tool-allowance figures were previously implied, not stated, in the scenario table below.]

Scenario	Monthly revenue	Monthly delivery hours	Contribution after stated allowances
3 clients, one of each offer	\$4,250	18	\$1,650
6 clients, two of each	\$8,500	36	\$3,800
12 clients, four of each	\$17,000	72	\$8,100
3 clients, twice the hours	\$4,250	36	\$300
3 clients, labor at \$125/hour	\$4,250	18	\$750

The first three setups book \$5,750, require 39 assumed delivery hours, and leave \$1,975 before fixed setup overhead under the base assumptions. At double setup hours, that contribution becomes **negative \$950**. Setup fees must fund setup work; recurring margin cannot hide loss-making onboarding.

Monthly tool allowances per client are \$75 Visibility, \$150 Lead Operations and \$100 Campaign, totaling \$325 for one of each. Contribution = revenue minus loaded delivery labor, tools, 10% sales allowance and \$500 fixed allowance. These scenarios do not predict customer acquisition or retention. Twelve clients is a staffing scenario, not a 90-day promise. At that scale the fixed-cost allowance may rise, so rerun it with actual staffing and tooling before approving expansion.

For the first cohort, propose reserving 15 hours/week from Dillon: 10 delivery/review and 5 discovery, documentation and coordination. Existing work commitments must be explicitly reconciled. Three setups need roughly four weeks of the 10-hour delivery allocation even before recurring service begins. Therefore **sell and schedule next week; do not promise all three implementations next week**. Start at most one new setup per week, never exceed two active builds, and calculate weekly capacity as setup hours plus recurring 6-hour-per-client budgets plus support/review. If projected weekly load exceeds 10 delivery/review hours, delay the next setup or reduce the recurring budget in writing. [Correction applied 2026-09-07 per QWEN27-REVIEW.md defect 2.] Reserve recurring capacity at 6 hours per accepted client per month (approximately 1.39 hours/week); support and review are included, not added twice. At three recurring clients, only about 5.8 of the 10 weekly delivery hours remain for new setup. Delay a start whenever setup plus existing commitments exceeds capacity. Delegate production only when an exact acceptance contract exists.

Sales motion and the next seven working days

- 1. Charter and terms.** Mac and Dillon settle the one-page mandate, offer scope, compensation mechanism, spending allowance and protected time. Decide who can approve which changes. Meeting output is a decision sheet, not just enthusiasm.
- 2. Choose the problem and buyers.** Mac/Jesse nominate five relationship-approved accounts. Run 20-minute discovery on current process, frequency, cost of failure, baseline, systems and owner. These counts are activity goals, not forecast conversion rates.
- 3. Build one proof.** Demonstrate a synthetic lead entering, being deduplicated, routed and drafted, then show an exception. Use the tested CSV and metadata components where they apply. Avoid a simulated dashboard labeled as live.

4. **Package the offer.** Adapt the existing editable Momentum deck, one-page scope, intake checklist and acceptance scorecard. Include proposed prices, setup schedule and support limit. Resolve the pending case-study/pricing decisions before sending externally.
5. **Conduct two qualified demonstrations.** Jesse handles agreed sales routing; Dillon diagnoses the workflow and implementation fit; Mac supports commercial decisions. Record why buyers accept or decline. Send proposals only through approved recipient routes.
6. **Onboard the first signed pilot.** Verify account owner, inputs, baseline and access. Implement in an isolated workspace; dry-run before live changes; record rollback and exception handling. Book additional starts according to capacity.
7. **Review and schedule.** Reconcile actual hours, spend, approvals and booked versus collected amounts. Improve the runbook from the first defects. Decide whether the next buyer is ready, not whether the models produced enough text.

Thirty, sixty and ninety days

Day 30: earn repeatability. Aim for one signed paid pilot in acceptance, two signed paid pilots scheduled, and a third signed or in written proposal. Count a pilot as signed only after proposal acceptance, access approval, first invoice or deposit, and a named operator. Do not count discovery calls or qualified demos as paid pilots. [Correction applied 2026-09-07 per QWEN27-REVIEW.md defect 5.] Track proposal accepted, contract signed, deposit collected, access ready and implementation accepted as separate milestones. A paid pilot requires collected payment; an invoice or discovery call alone does not qualify. Advance only after acceptance evidence, a usable operator handoff, time logs and a recurring cost estimate exist. Require zero known critical routing or data-loss defects. A 60% gross-margin planning target is a decision threshold, not an industry benchmark. If actual gross margin is below 50%, reprice or reduce scope before adding volume.

Day 60: add measured capacity. Consider six active clients only after two consecutive delivery cycles meet their promised dates, scope and review budget. Record reasons for every late task and rework cycle. Bring in a part-time production/QA specialist when forward four-week demand exceeds 70% of the approved delivery capacity. The 30% reserve is a proposed buffer for exceptions. A trained reviewer must be able to reproduce the acceptance check without Dillon narrating every step.

Day 90: choose a scale path. If retention, collections, referenceable outcomes and contribution support it, standardize the most profitable repeated workflow and staff toward the twelve-client scenario. Otherwise narrow to the strongest offer. Consider a reusable software product only after several clients repeatedly buy essentially the same process; custom features for one client are not yet a product market.

Ownership and compensation discussion

Proposed responsibilities: Mac owns strategic sponsorship, final pricing and commercial authority; Dillon owns solution design, delivery system and quality; Jesse owns agreed prospecting and sales follow-through; the existing account owner owns client communication continuity; a named operator handles exceptions and support. Codex remains Dillon's orchestrator and verifier; Hermes and other models execute bounded tasks.

Protect Dillon's time before accepting a commission-only expansion. Preferred discussion structure: a defined paid capacity allocation or role adjustment, plus a transparent incentive on collected contribution or revenue, with sales credit separately defined. If a temporary commission-only pilot is chosen, cap hours and duration and specify expense reimbursement. Do not accept an undefined share of “net profit” with unbounded overhead allocation.

Put the following in writing: eligible revenue, sales versus delivery credit, renewal ownership, collection trigger, refunds/nonpayment, tool reimbursement, support duty, pre-existing versus newly built IP, and the review date. Percentages need Mac's agreement and a reconciled model. Qwen's suggested percentages were not sourced and were excluded.

The operating system and model roles

Use the canonical client registry and queue already present. Resolve-Client.ps1 routes identity; existing Marketing Chief execution graphs define work; Accept-WorkerHandoff.ps1 supports reviewed reconciliation. Do not build a second client database because a new harness has memory.

Each job receives a compact source packet: exact client/brand, objective, timestamped evidence, source locations, permitted tools, output path/schema, acceptance check, time/token cap and stop condition. Retrieve relevant context incrementally. A full Slack archive in every prompt is costly, mixes accounts and does not provide durable memory.

Small Qwen: extraction, normalization, repetitive first drafts and narrow code. Larger Qwen: synthesis, strategy alternatives and difficult implementation drafts. Deterministic Python: counts, economics, schema checks and reconciliation. Codex: current research, connector retrieval, decisions, review and final action verification. DeepSeek Harness is a candidate execution shell; installing it does not add model intelligence or automatically grant connectors.

Run plan □ critique □ repair □ acceptance as bounded stages. Log task success, critical errors, human repair minutes, cost per accepted artifact and latency. GPU tokens/second measures generation throughput; it does not establish sales capacity or dependable autonomy. Tonight's unmodified coding drafts had defects; Codex repair remains part of the real cost.

Forecasting starts with instrumented data. Initially use transparent scenarios and due-date/rule-based alerts. After sufficient comparable history exists, backtest a predictor against a simple baseline and report intervals and error. Do not market a “predictor of all business problems.”

Risks and concrete responses

Scope creep: one workflow/domain/brand per offer, written change request and rescheduling. Client access delays: verify feasibility before committing dates. Weak adoption: name an operator, train them and track real usage. Incorrect output: independent checks and source readback. Support overload: capped workload and paid escalation. Margin illusion: include review, rework and founder time. Demand risk: obtain paid pilot commitments before hiring. Commodity pricing: sell verified implementation and support. Tool churn: pin versions and change only after the same acceptance packet improves.

The market already contains implementation specialists; Zapier's directory demonstrates an established services category rather than an empty market. Momentum should compete on segment knowledge and delivery evidence. [Zapier partner directory](#). For architecture, begin with a simple workflow and add agents when the task requires them; this follows the practical implementation guidance rather than maximizing agent count. [OpenAI agent guide](#).

Meeting decision sheet

- First customer segment and the three founding accounts, with relationship owners.
- Flagship offer, exact scope/prices/terms and exclusions.
- Dillon's protected hours, role authority, compensation and tool reimbursement.
- Jesse's sales role and the exact commission formula.
- Client-facing claims, case-study permission and the existing deck's pending decisions.
- Named support owner, delivery calendar and a 30-day review date.
- A small operating budget and the conditions that permit increasing it.

The recommendation is ready for discussion. Commercial approval and customer demand remain to be established; the tools and model experiments support delivery but cannot substitute for either.

Extension — 2026-09-07: the division already has a machine

Appended 2026-09-07. This section extends the September 4 proposal; it does not replace it. Everything above still stands, including its corrections. Where this extension contradicts an earlier statement, it says so explicitly and gives the measurement.

Vault evidence for this section: [\[\[12_Brain/05_Projects/2026-09-07 - Momentum AI division launch\]\]](#).

The thesis, and whether it survives contact with the evidence

Dillon stated the product to Mac Frederick in #ai-tech-news:

"we can create skills inside of Claude basically organizational skills and you'll just tell the code who the client is and you'll basically build in all of these pre-existing things that would normally come in like an overview deck for example and it really automate that process while keeping a constant brand guideline continuity across all PowerPoint that never really like fails the design test."

Read strictly, that is a claim with three parts: (1) a per-client context switch, (2) reusable organizational skills that carry Momentum's pre-existing assets, (3) brand continuity that does not fail the design test. Tested against what is on disk on 2026-09-07:

Claim	State	Evidence
Per-client context switch	Holds. A momentum-client-context skill resolves the named client, and the canonical registry plus Resolve-Client.ps1 already route identity.	Skill roster; registry at client-operations/registry/clients.json.
Reusable organizational skills	Holds in substance, fails in provenance. Five Momentum skills resolve at the account layer — momentum-client-context, momentum-brand-system, momentum-client-intake, momentum-client-report, momentum-spec-homepage. None has a locatable SKILL.md anywhere under C:\Users\dillo\.claude.	Filesystem search 2026-09-07 found no SKILL.md under any momentum-* path.
"Never fails the design test"	Holds, and is now measurable rather than rhetorical. The design system carries 31 measured contrast pairs with 0 failures, and its build script refuses to render if any pair fails.	2026-09-05-momentum-design-system/audit/contrast.json, build.py.

The thesis holds. The correct framing for the division is not "we sell AI services." It is: *the delivery machinery exists, it is measured, and the measurement is the thing competitors cannot copy from a pricing page*. Two qualifications belong in any client-facing use of it:

1. "Never fails the design test" is true of the **new system**, not of what Momentum has already shipped. The same audit measured 120 token pairs across ten live prospect builds and found **37 AA failures**. The claim is a forward promise backed by a working gate, not a description of the existing estate.
2. The five skills are the machine, and the machine's source is not in a repository Dillon controls. That is a productisation risk, not a delivery risk — the skills work today. See "Open items" below.

Why this changes the plan's centre of gravity

The September 4 plan sells three scoped offers. It is correct and should ship. What it does not yet do is state *why a buyer should pick Momentum over the implementation specialists the plan itself concedes already exist* (Zapier's partner directory). The answer is now measurable and it is this:

Momentum can show which actual humans a client's conversions were. Almost no SMB agency can.

That is not a positioning line invented for a deck. It is the direct consequence of an operating defect found this week, and it is worth more as a differentiator than as an embarrassment. See "Match-back" below.

Market position — anchors for the pricing decision

Sourced from a scan run in the AI division session, **not from primary vendor pages**. Treat as a planning anchor with a verification gate, not as a quotable competitive claim. Full note with expiry:

[[12_Brain/06_Research/2026-09-07 - AI search retainer pricing and Philadelphia position]].

- AEO retainers in the SMB tier cluster at **\$1,500–3,000/month** entry and **\$3,000–8,000/month** for active work.
- **WebFX (Harrisburg)** is the only Pennsylvania agency found publishing an AI-search price, at **\$3,000/month**.
- **ChatGPT Ads went self-serve** for all US businesses on **2026-05-05** — cost-per-click, no minimum.
- **No Philadelphia SMB agency has launched a named "AI division."**

The fourth item is the positioning, not a footnote. Momentum's advantage window is the gap between now and the first Philadelphia competitor naming a division, and that window is measured in months.

Consequence for the offer table above. The AI Visibility Program is priced at \$1,500 setup / \$1,500 monthly. That sits at the *bottom* of the scanned entry band and 50% under the only published Pennsylvania anchor. Two readings, and Mac should pick one deliberately rather than by drift: it is a founding-cohort price that rises after three references, or it is underpriced. It should not be both.

Mac's open asks for the sales deck

Mac has four open asks: **AEO, AI Marketing, Social Packages, pricing**. needmomentum.com has no AEO page, so scope and price are drafted here and **flagged for his sign-off**. Nothing below is approved or sendable.

AEO — draft scope and price for sign-off

Reuse the AI Visibility Program scope defined above verbatim; do not author a second competing scope. The open decisions are price positioning and the named deliverable.

Decision	Draft	Anchor
Entry price	\$1,500 setup / \$1,500 monthly, founding cohort only, first three clients	Bottom of scanned \$1,500–3,000 entry band
Standard price after three references	\$2,500/month	Mid-band; still under WebFX's \$3,000
Named deliverable	The 20-question baseline across two engines, re-measured monthly with engine, question, date and source on every observation	The audit trail is the product; a ranking number is not
Guarantee language	None. Observed presence is reported; rankings, citations, leads and revenue are explicitly not guaranteed	Already settled in the offer table; do not reopen under sales pressure

AI Marketing and Social Packages

These are Mac's navigation categories, not new delivery organizations — that decision from the September 4 plan stands. AI Marketing maps to the Campaign Production System (\$750 / \$1,250). Social Packages have **no scope, no price and no delivery contract drafted**, and should not be listed on a deck until they do. Listing a category with no acceptance contract behind it is how the scope creep in the risk register actually happens.

ChatGPT Ads

Self-serve since 2026-05-05 with no minimum removes the entry barrier that made this a specialist service. The existing signed \$500/month management lane is the price anchor Momentum already has. This stays a separate service decision, as stated above.

Match-back: the differentiator, and the defect underneath it

Conversion-to-named-lead match-back — telling a client which actual people their reported conversions were — has been promised **nine times to Omega Landscape, twice to Kimberly James Bridal, twice to Onsite Construction**, and currently **blocks the Fresh Blends deck**.

It cannot be produced. The Zapier notifications arriving in Gmail carry no lead data — only a link reading "open this Zap run in Zapier." **Fagan Painting is the only account where the loop closes**, and the reason is narrow and fixable: its leads are parsed into the email body.

The fix is a Zap configuration change, not a reporting effort. That distinction is the whole finding. Thirteen promises across four accounts have been treated as a reporting backlog; they are one integration defect with a known-good reference implementation already running on a live account.

Consequences for this plan:

1. **Fixing it is the Lead Operations Pilot's own proof.** The flagship offer's acceptance criteria include "destination readback succeeds for 100% of accepted test records." Momentum currently fails that criterion on its own book of business. Fix it internally first, and the pilot demo stops being synthetic — it becomes a case study with named clients and dates.
2. **It is the sharpest sales claim available.** "We tell you which actual humans your conversions were" is concrete, verifiable, and outside the reach of an agency whose reporting stops at platform totals.
3. **It must not be sold before it is fixed.** The claim is currently true for one account out of four. Selling it now would repeat the thirteen promises at a larger scale.

Supporting vault note: [[12_Brain/03_Concepts/2026-09-07 - Conversion match-back is the differentiator]]. Independent corroboration of the access dependency: Omega's Google Ads lead export through Zapier and Sheets could not be authorized — Dillon holds Standard access, elitelandscapingconcrete@gmail.com is the account Admin (01_Clients/Omega Landscaping/Agent Memory.md).

The design proof, with the numbers corrected

The design system makes "never fails the design test" a build gate rather than a claim. The mechanism is worth stating because it is the reusable part: **no single Momentum orange clears 4.5:1 as text on a light ground while also working as a button fill**. #f58320 scores 7.19:1 as a field against near-black and 2.58:1 as text on white. One token cannot do both jobs, so the accent splits by job.

Corrections to the figures circulating in the AI division session:

Circulated	Measured	Source
Tokens named --m-signal / --m-signal-ink	--m-accent (field), --m-on-accent (label on field, 7.19:1), --m-accent-ink (text on paper, 6.18:1). No signal token exists.	tokens/momentum.tokens.css — 0 occurrences of "signal"
29 pairs measured, 0 failures	31 pairs, 0 failures. 22 are small-text pairs.	audit/contrast.json
Lowest text pair 4.695:1	Lowest small-text pair 4.54:1 ("nav mark if nav ever sits on brand"). Lowest pair of any size is 3.14:1 , declared large-text-only and stated rather than hidden.	audit/contrast.json
needmomentum.com: 4 type families, 15 font sizes, 10+ radii	5 font families (Roboto, Open Sans, Poppins, Hind, bare sans-serif), 20 distinct radii , 14 transition pairs. 2 oranges on one page. Body copy #888888 on #f5f5f5 = 3.25:1 .	AUDIT.md §2.4, measured in-browser at 1440×900
10 radar sites, QA sheet recorded 10/10 PASS	Not located as stated. The 10/10 figure found on disk is the garage-door missed-call proof-of-concept scenario pass rate, a different artifact. The radar QA receipt that was found (2026-09-05, 20 sites) reports a DEGRADED detector bill, not a clean pass.	radar-2026-09-05-build-report.md

The drift itself is client-affecting and is the honest "before" half of the story: **Nolt's Auto Parts and Weathers Motors both shipped the kit's unmodified demo palette** (--accent: #d4762a on --brand: #1d4e6d), and advance-exterior carries the same demo navy. --muted on --panel fails AA on **9 of 10** sites at 4.28–4.49:1 — nine consecutive near-misses, which is a derivation formula validating against a surface the text never appears on, not ten separate mistakes. One hero eyebrow renders at **1.01:1**, which is invisible.

This is a sellable narrative, not a liability, provided it is told in the right order: we measured our own work, found 37 failures, built a gate that makes them impossible, and the gate ships with every client build. An agency that can say that is making a verifiable claim. See [[12_Brain/03_Concepts/2026-09-07 - One token cannot do two jobs]].

Control-system gaps that gate the launch

The plan assumes a working operating system underneath the offers. Several parts of it are not currently observable. These are launch dependencies, not housekeeping.

- **Five clients have live Slack channels and no vault record, Nexla among them** — and Nexla is the only client with **live ad spend invisible to the control system**. A division that sells "evidence of what changed" cannot have a spending account outside its own evidence layer.
- **Vault folders diverge from Slack channels**: Omega Landscaping vs #omega-landscape, Onsite Concrete vs #onsite-construction. A Shadow HVAC folder holds a complete Next.js build with no Slack channel at all.
- **slackChannels is populated for 5 of 24 registry records**. The registry is the routing authority the plan depends on in "The operating system and model roles."

- **Inbox, 2026-09-04: 104,403 messages, 93,334 unread, 181 drafts — 57 of them unsent replies to paying clients**, none older than 38 days. The Client Triage, Client Triage/Drafted and KJB Daily Leads labels are all empty, so the filters either never fired or stopped.
- **Align HCM is stale in the registry** — listed active with six deliverables including Dayforce exhibitor artwork with a rush fee, while CLAUDE.md line 16 and System/operating-status.md both record it ended, confirmed 2026-09-02. Closed out at [[12_Brain/04_Decisions/2026-09-07 - Align HCM registry record is superseded]].

Full pass: [[12_Brain/07_Reviews/2026-09-07 - AI division evidence pass]].

Two client threads that bear on the launch

Puttery — zero of seven access gates cleared as of 2026-09-04. Vendor eligibility **is** confirmed (Business Group 28086, Business 37824). Joe Pedevillano asked on 2026-08-25 to confirm click-to-reservation matching and has had no reply. Momentum side: Mac Frederick, Melissa Silber, Jesse DiLaura. "Lauren" is **Laura at Resy**, api-integration@resy.com, four items outstanding. This confirms the September 4 judgment that Puttery is an engineering proof opportunity, not the simplest first rollout.

Bridge Software Development — zero dated client acceptances against three paid milestones; Milestone 3 is complete but sits in an unmerged PR. **Correction, recorded explicitly:** an earlier analysis attributed missed reviews on 18 August, 23 August and 1 September to Dillon. That was wrong. Those were the client's slips, and the 1 September slot was declined by Tori. **Momentum is the fastest-moving party on Bridge.** This matters to the plan because Bridge is one of the four registry-affiliated accounts named in "Initial customer and qualification," and a false record of Momentum lateness would have argued against using it.

Security — exposure state, locators only

No credential values appear here or in any vault note from this pass.

- A live Tock credential sat in plaintext in Gmail thread 1a049e62f063f279, sent 2026-08-31 to four recipients. A revocation was drafted 2026-09-02 and **never sent** — it remains queued in System/approval-queue.md.
- Four further sets of plaintext credentials were found in Slack; the oldest has been exposed 24 days.
- `_os/test/public-safety.test.js` must continue to pass on every vault change.

A division whose pitch is "we connect your systems and verify the result" cannot carry an unrotated credential exposure into its first client conversation. Rotation is approval-gated and sits in the queue.

Open items — none of these are blockers, all of them are decisions

1. **Skill provenance.** The five Momentum skills have no SKILL.md on this machine. If they are the delivery machinery, their source belongs in a repository Dillon controls, with the same treatment the design system gets. Until then the machine is real but not portable, and not auditable.
2. **Skill collision.** client-report lives at dillon-os/.claude/skills/ (28 skills there, not the eight previously believed) and renders branded HTML into Daily-Briefs/reports/. momentum-client-report lives at the account layer and leads with match-back. Both trigger on "client report", "monthly report", "performance recap". **Nothing is deleted.** The decision is which one owns that phrase, and the answer is probably that the vault skill renders and the Momentum skill supplies the match-back spine — but that has to be written down, not left to dispatch order.
3. **Social Packages** need a scope and an acceptance contract before appearing on any deck.
4. **AEO price positioning** — founding-cohort rate or standard rate. Mac's call.
5. **The match-back Zap fix** must land before match-back is sold.
6. **Nexla** must become visible to the control system before the division claims an evidence layer.

What this extension does not change

The three offers, their acceptance contracts, the economics, the capacity arithmetic and the seven-day sales motion above are unchanged. The 15 hours/week reservation, the one-setup-per-week ceiling and the two-active-builds ceiling still bind — and the work catalogued in this extension is exactly the kind of internal work that consumes them. None of the fixes listed here are inside a client's six-hour monthly budget.

Corrections, 2026-09-07 evening

Two statements above are sharpened by the evening's research pass
(2026-09-07-google-aistudio-batch/research/aeo-geo-topic-research.md):

1. **WebFX's \$3,000/month is now verified** on WebFX's own GEO page, and its May 2026 cost article brackets small-business agency GEO at \$1,500–\$5,000/mo. Market anchor 2 moves from "scan" to "primary source"; anchor 1 is bracketed by the same source.
2. **"AEO has no page on needmomentum.com" is true of a service page only.** A blog post exists — needmomentum.com/answer-engine-optimization/, 2026-02-08, "The Rise of Answer Engine Optimization", Momentum Mondays — with no price, package or client. The AEO offer still needs its own page; the article is a link, not a substitute.

Also from that pass: demand is on the word **AEO** ("what is aeo", 4,400/mo US, growing from 140 two years ago), not "GEO" (declining). Name the offer for the term that is growing. needmomentum.com is cited 81 times in Google AI Overviews — fifth of six named regional agencies — and none of the sampled citations is on an AI topic. That is the gap the AI Visibility Program exists to close, on Momentum's own site first.