

House stack vs Gemini 3.8 Flash

Dillon Mohr · Momentum Digital

PRIVATE REVIEW EDITION · SEPTEMBER 2026

Complete source draft, formatted for your review. Publication status and open decisions are preserved inside.

Before this goes public

This is an internal review copy. Offers, pricing, first-person claims, client examples and delivery promises require their existing approvals before public use. This PDF export does not change the source approval status.

Source: PRICE.md

Typography: Archivo Black display, Georgia reading text and Calibri utility text. Georgia is a print-specific reading choice; the existing Momentum web design system is unchanged.

Inside

01 / The one number

02 / API list: how many times more than 3.8 intro

03 / Same work, cash on the meter

04 / Seats are a different meter

05 / What this means for Mac week

House stack vs Gemini 3.8 Flash

Read 2026-09-07. API list is USD per 1 million tokens. Seats are monthly prepaid. Recheck live pages before a spend decision.

Short answer: Gemini 3.8 Flash intro list is **2.7× cheaper than Sonnet 5**, **6.7× cheaper than Opus 5**, and **13× cheaper than Fable 5 and GPT-6 Astra** on both input and output. It is **not** the cheapest thing in the house. GPT-5.6 Luna and Cursor Composer 2.5 undercut it.

Benchmarks for the same set live in BENCH.md. Headline: BenchLM #6 at **78.4**, ties Opus on DeepSWE, dies on Terminal-Bench 4.0 (**19.1** vs Opus **52** / Astra **58**). Luna max is close on TB 2.1, not on the composite. Composer has no Sep 2026 frontier card.

The one number

Gemini 3.8 Flash (gemini-3.8-flash) intro through **2026-12-31**:

	Input	Output	Cache read
Google list	\$0.75	\$3.75	\$0.075
Cursor Other Models	\$0.75	\$3.50	\$0.075

Standard list from **2027-01-01: \$1.50 / \$7.50**. Thinking tokens bill as output. Google says 3.8 "works harder" and often emits ~30% more output tokens than a same-task peer, so real savings vs Sonnet land nearer **2.0×**, not 2.7×, on long agent turns.

API list: how many times more than 3.8 intro

Model	Input	Output	vs 3.8 input	vs 3.8 output	Notes
Gemini 3.8 Flash	\$0.75	\$3.75	1.0×	1.0×	Intro through Dec 31
GPT-5.6 Luna	\$0.20	\$1.20	0.27× (Luna cheaper)	0.32×	Codex cheap tier. Official OpenAI list.
Composer 2.5	\$0.50	\$2.50	0.67× (Composer cheaper)	0.67×	Cursor first-party pool
Haiku 4.5	\$1	\$5	1.3×	1.3×	Claude grunt
Sonnet 5	\$2	\$10	2.7×	2.7×	Permanent \$2/\$10 as of Aug 10
Grok 4.6	\$2	\$6	2.7×	1.6×	Doubles at ≥200k prompt
GPT-5.6 Terra	\$2	\$12	2.7×	3.2×	Codex mid
Sonnet 4.6	\$3	\$15	4.0×	4.0×	Friday fallback
GPT-5.6 Sol (promo)	\$4	\$20	5.3×	5.3×	Promo at least through Nov 21
Grok 4.6 Fast	\$4	\$12	5.3×	3.2×	This cloud run's family
Opus 5	\$5	\$25	6.7×	6.7×	Fast mode 2× price
Fable 5 / 5.1	\$10	\$50	13×	13×	Same list
GPT-6 Astra	\$10	\$50	13×	13×	Whole request → \$20/\$75 if input >272k. Fast 2×.

Percent cheaper on list (1 – 3.8 / other), intro rate, ignoring the 30% token bump:

- vs **Sonnet 5: 62% cheaper**
- vs **Opus 5: 85% cheaper**
- vs **Fable 5 / Astra: 92% cheaper**
- vs **Grok 4.6 output: 38% cheaper** (input still 62%)
- vs **Sol promo: 81% cheaper**

After Jan 1, 3.8 at \$1.50 / \$7.50 is still **25% cheaper than Sonnet 5** and **70% cheaper than Opus 5**.

Same work, cash on the meter

Blend used below: **15M input / 6M output**, no cache, standard (not Fast), short context. That is a heavy Codex-style week, not an invoice.

Model	Week list	vs Gemini \$33.75	With 3.8's +30% output (\$40.50)
Luna	\$10	Gemini costs 3.3× more	—
Composer 2.5	\$23	Gemini costs 1.5× more	—
Gemini 3.8 Flash	\$34	—	\$41
Grok 4.6	\$66	2.0×	1.6×
Sonnet 5	\$90	2.7×	2.2×
Terra	\$102	3.0×	2.5×
Grok 4.6 Fast	\$132	3.9×	3.3×
Sol promo	\$180	5.3×	4.4×
Opus 5	\$225	6.7×	5.6×
Fable / Astra	\$450	13×	11×
Astra Fast	\$900	27×	22×
Astra >272k	\$750	22×	19×

Two Astra weeks at that blend: **~\$900 list vs ~\$68–\$81 Gemini**. That is the weekend-overflow math. It is **not** cash returned if those weeks were included Codex banks or compensation resets.

One Design-flush turn (20k in / 8k out, no cache):

Model	Turn
Luna	\$0.014
Composer 2.5	\$0.030
Gemini 3.8	\$0.045 (\$0.054 at +30% out)
Grok 4.6	\$0.088
Sonnet 5	\$0.120
Opus 5	\$0.300
Fable / Astra	\$0.600

Seats are a different meter

API ratios do not refund a prepaid seat. Moving overflow to Gemini only saves money when the alternative is **credits, a weekly reset, on-demand Cursor Other Models, or a new Ultra/Max/Pro upgrade.**

Seat	Monthly	What it buys	Where Gemini sits
Claude Pro	\$20 (\$17 annual)	Claude Code + Design. Fable is usage-credits only.	Not inside this seat
Claude Max 5x / 20x	\$100 / \$200	5x or 20x Pro per 5-hour session. Same models.	Not inside this seat
ChatGPT Plus	\$20	Codex + GPT-5.6 family. Shared 5-hour + weekly bank.	Not inside this seat
ChatGPT Pro 5x / 20x	\$100 / \$200	5x or 20x Plus Codex. Astra burns this bank fast.	Not inside this seat
Cursor Pro / Pro+ / Ultra	\$20 / \$60 / \$200	Cursor Models (Grok 4.6, 4.5, Composer 2.5) + Other Models included \$20 / \$70 / \$400	3.8 in Cursor is Other Models at \$0.75 / \$3.50
Google AI Pro / Ultra 5x / 20x	\$20 / \$100 / \$200	Consumer Gemini + Flow / Antigravity headroom. Not an API refund.	CLI Google-login for individuals is a separate, flaky rail. API key is the clean cheap path.

Cursor docs: Pro / Pro+ / Ultra include two pools. Grok and Composer draw the generous first-party pool. Claude, Gemini, and OpenAI draw Other Models at API list. This cloud run (cursor-grok-4.6-xhigh-fast) is first-party **Fast** (\$4 / \$12), not Gemini.

Do **not** switch a Cursor Grok session to Gemini 3.8 to "save money." That moves work from the included Cursor Models pool onto the small Other Models dollar cap.

What this means for Mac week

1. **Coding overflow after Codex dies:** Gemini 3.8 Flash API at \$0.75 / \$3.75. That is the 13x Astra hedge.
2. **Mac / Snapshot / GDPval writing:** stay on Claude Code Sonnet 5. Saving ~\$0.075 a Design turn is not worth the knowledge-work gap (GDPval 1545 vs Sonnet 1618 / Opus 1861).
3. **Hard agent closer (TB 4.0-class):** Astra if the bank is there, else Opus 5. Not Gemini (official TB 4.0 19.1 vs Opus 51.8).
4. **Inside Cursor:** keep Grok 4.6 / Composer 2.5. Composer is cheaper than 3.8 on list and sits in the included pool.
5. **Do not buy** Claude Max, ChatGPT Pro, Cursor Ultra, or Google Ultra this week to chase 3.8. Buy a Gemini API key only if Google-login 429s.

Sources: [Google Gemini 3.8 Flash](#), [Google Agent Platform pricing](#), [Anthropic pricing](#), [OpenAI API pricing](#), [xAI pricing](#), [Cursor models and pricing](#), [Cursor plan help](#).